

I found that the further I deviated from 1 with the x value, the funnier my result was. I began with the P + 7 portion of the assignment before trying any other inputs, and the results were somewhat interesting, but lacked any kind of flavor. When I went to numbers lower than 7, my results were just trying to make the endings of the sentences as basic as possible. After P + 10 i started jumping up by 10s, and by the time I got to 100, I was enjoying my results a lot more. I landed on P + 300, which gave me a poetic ending to each sentence, creating a new piece of writing with a new flavor. It feels very far removed from the original poem despite only changing the last word. To apply the P + 7 technique to every noun in the sentence, I would make it so that whenever the pipeline reaches a token that has been identified as a noun, you take the model's ranked probability distribution for that position and choose the seventh-highest item. That selection replaces the original token while everything else in the text passes through unchanged, and the process repeats independently for each noun in the input.